

pavo: Perceptual Analysis, Visualization and Organization of Spectral Color Data in R

Rafael Maia¹, Paul-Pierre Bitton², and Chad Eliason¹

¹Integrated Bioscience PhD Program, University of Akron, Akron OH

²Department of Biological Sciences, University of Windsor, Windsor ON

November 29, 2012

Contents

1 Dataset Description	1
2 Organizing and Processing Spectral Data	2
2.1 Importing Data	2
2.2 Processing Data	3
2.2.1 Averaging Spectra	3
3 Visualizing Spectral Data	6
4 Analyzing Spectral Data	7
4.1 Overview	7
4.2 Spectral Shape Analysis	8
4.2.1 Trichromatic variables	8
4.2.2 Peak shape descriptors	10

Introduction

pavo is an R package developed with the goal of establishing a flexible and integrated workflow for working with spectral color data. It includes functions that take advantage of new data classes in order to work seamlessly from importing raw data to visualization and analysis.

Although **pavo** deals largely with spectral reflectance data from bird feathers, it is meant to be applicable for a range of taxa and applications. It provides flexible ways to input spectral data from a variety of equipment manufacturers, process these data, extract variables and produce publication-quality graphics.

pavo was written with the following workflow in mind:

1. **Organize** spectral data by inputting files, processing spectra (e.g., to remove noise, negative values, smooth curves, etc.)
2. **Analyze** the resulting files, either using typical tristimulus color variables (hue, saturation, brightness) or using visual models based on perceptual data from the taxon of interest.
3. **Visualize** the output, with multiple options provided for exploratory analyses.

Below we will show the main functions in the package in an example workflow.

1 Dataset Description

The data used in this example is available from [github by clicking here](#). You can download and extract it to follow the vignette.

The data consists of reflectance spectra obtained using Avantes equipment and software from seven bird species: Northern Cardinal (*Cardinalis cardinalis*), Wattled Jacana (*Jacana jacana*), Baltimore Oriole (*Icterus galbula*), Peach-fronted Parakeet (*Aratinga aurea*), American Robin (*Turdus migratorius*), and Sayaca Tanager (*Thraupis sayaca*). Several individuals were measured (sample size varies by species), and 3 spectra were collected from each individual.

The samples do not have the same sample sizes and have additional peculiarities that should emphasize the flexibility `pavo` offers, as we'll see below.

2 Organizing and Processing Spectral Data

2.1 Importing Data

The first thing we need to do is import the spectral data into R using the function `getspec()`. Since the spectra were obtained using Avantes software, we will need to specify that the files have the ".ttt" extension. Further, the data is organized in subdirectories for each species. `getspec` does recursive sampling, and may include the names of the subdirectories in the spectra name if desired. A final issue with the data is that it was collected using a computer with international numbering input, which means it uses commas instead of periods as a decimal separator. We can specify that in the function call.

I have downloaded the file and placed it in a directory called `~/github/pavo/vignette_data`. By default, `getspec` will search for files in the default folder, but a different one can be specified:

```
> specs <- getspec("~/github/pavo/vignette_data/", ext="ttt", decimal=",",
+                 subdir=T, subdir.names=F)
> specs[1:10,1:4]

      w1 cardinal.0001 cardinal.0002 cardinal.0003
1  300      5.7453      8.0612      8.0723
```

2	301	6.0181	8.3926	8.8669
3	302	5.9820	8.8280	9.0680
4	303	6.2916	8.7621	8.7877
5	304	6.6277	8.6819	9.3450
6	305	6.3347	9.6016	9.4834
7	306	6.3189	9.5712	9.3533
8	307	6.7951	9.4650	9.9492
9	308	7.0758	9.4677	9.8587
10	309	7.2126	10.6172	10.5396

```
> dim(specs) #data has 214 spectra, from 300 to 700 nm
```

```
[1] 401 214
```

When `pavo` imports spectra, it creates an object of class `rspec`, which inherits attributes from the `data.frame` class:

```
> is.rspec(specs)
```

```
[1] TRUE
```

If you already have multiple spectra in a single data frame that you'd like to use with `pavo` functions, you can use the command `as.rspec` to convert it to an `rspec` object. The function will attempt to identify the wavelength variable or, if it doesn't have one, it can be specified in the function call.

2.2 Processing Data

2.2.1 Averaging Spectra

As previously described, our data constitutes of multiple individuals, and each was measured three times, as is common in order to avoid measurement bias. A good way to visualize the repeatability of our measurements is to plot the spectra of each individual separately. The function `explorespec` provides an easy way of doing so. You may specify the number of spectra to be plotted in the same panel using the argument `specreps`, and the function will adjust the number of panels per page accordingly. We will exemplify this function using only the 12 cardinal individuals measured:

```
> explorespec(specs[,1:37], specreps=3)
> # 36 spectra plus the first (wl) column
```

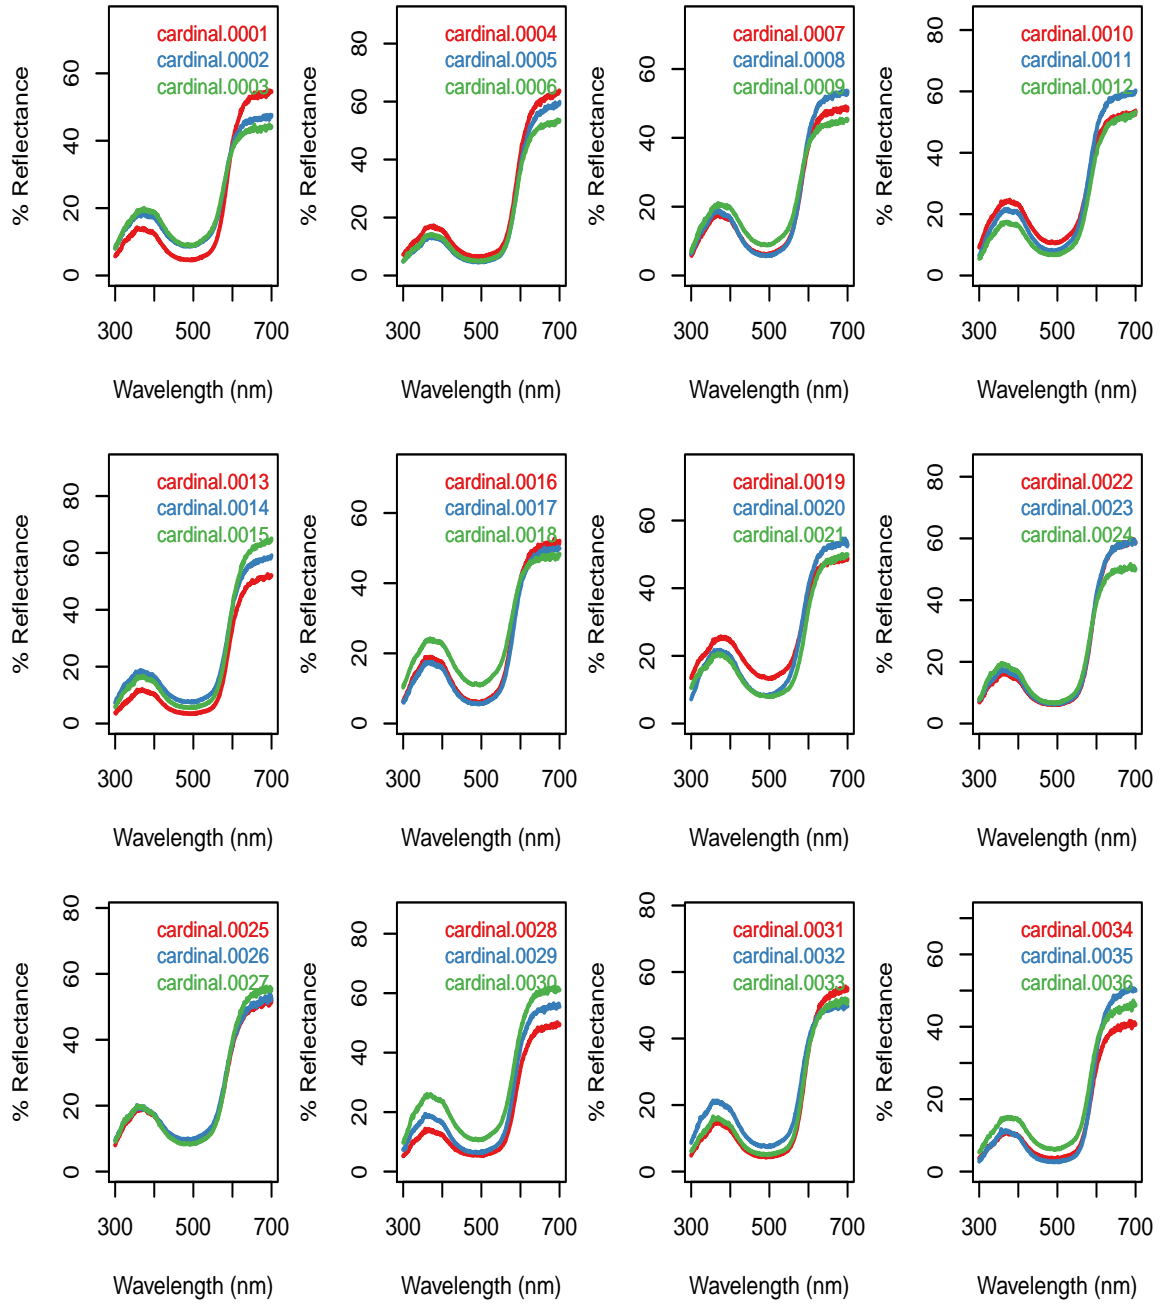


Figure 1: Result from `explorespec`, showing the three measurements for each individual in separate panels

So our first step would be to take the average of each of these three measurements in order to obtain average individual spectra to be used in further analyses. The function `aggspec` does this. The `by` argument can be either a number (specifying how many specs should be averaged for each new sample) or a vector specifying the identities of the spectra to be combined (see below):

```
> mspecs <- aggspec(specs, by=3, FXN=mean)
> mspecs[1:5, 1:4]
```

	wl	cardinal.0001	cardinal.0004	cardinal.0007
1	300	7.292933	5.676700	6.387233
2	301	7.759200	5.806700	6.698200
3	302	7.959333	5.858467	6.910500
4	303	7.947133	6.130267	7.357567
5	304	8.218200	6.127933	7.195267

```
> dim(mspecs) #data now has 72 spectra, one for each individual
```

```
[1] 401 72
```

Now we'll use the `aggspec` function again, but this time to take the average spectrum for each species. However, each species has a different number of samples, so we can't use the `by` argument as before. Instead we will use regular expressions to create a species name vector by removing the numbers that identify individual spectra:

```
> # create a vector with species identity names
> spp <- gsub('\\.[0-9].*$', '', names(mspecs))[-1]
> table(spp)
```

spp	cardinal	jacana	oriole	parakeet	robin	tanager
	12	9	9	13	10	18

Instead, we are going to use the `spp` vector we created to tell the `aggspec` function how to average the spectra in `mspec`:

```
> sppspec <- aggspec(mspecs, by=spp, FXN=mean)
> sppspec[1:5, ]
```

	wl	cardinal	jacana	oriole	parakeet	robin
1	300	7.049397	7.334781	3.889693	7.629954	3.981747
2	301	7.254161	7.354033	3.905322	7.746882	3.914297
3	302	7.444275	7.452556	4.126619	7.886877	4.187073
4	303	7.820686	8.085541	4.390685	8.491367	4.507410
5	304	7.843394	7.714526	4.183637	8.658000	4.068800
		tanager				
1		9.021043				

```

2 9.525854
3 9.405980
4 10.199843
5 9.684522

```

```
> explorespec(sppspec, 6)
```

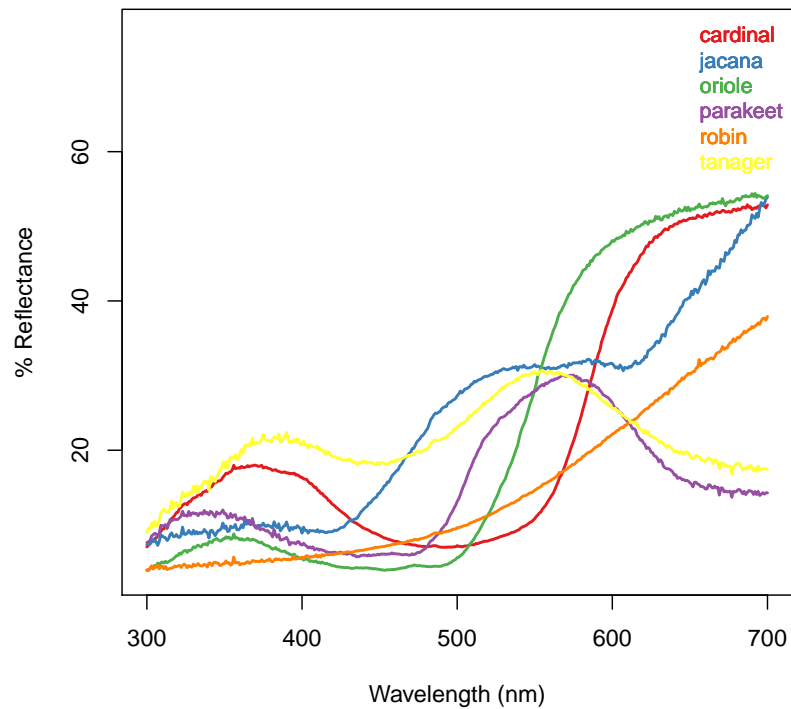


Figure 2: Result from `explorespec` for species means

3 Visualizing Spectral Data

```

> data(sicalis)
> par(mfrow=c(1,2))
> plot(sicalis, type='o', col=spec2rgb(sicalis))
> plot(procspec(sicalis, opt='smooth'), type='o', col=spec2rgb(sicalis))

```

processing options applied:
smoothing spectra with a span of 0.25

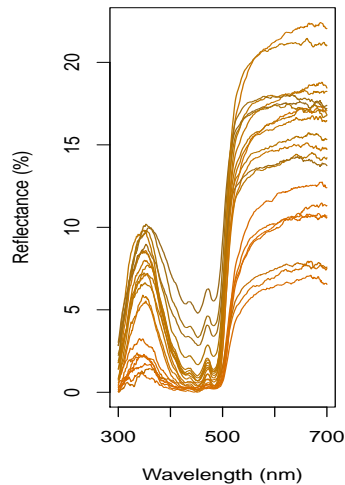


Figure 3: Overlay plot with colors calculated from human color matching functions

```
>
> #plot(sicalis, type='s', col=spec2rgb(sicalis))
```

Colors can be mapped to spectra using `spec2rgb` as shown in Figure 3.

4 Analyzing Spectral Data

4.1 Overview

`pavo` offers two main approaches for spectral data analysis. First, variables can be calculated based on the shape of the reflectance spectra. By using special R classes for spectra data frame objects, this can easily be done using the `summary` function in an `rspec` object (see below). The second function for spectral shape analysis is `peakshape`, which returns descriptors for peaks in spectral curves, as outlined below.

Second, reflectance spectra can be analyzed accounting for the visual system receiving the color signal, therefore representing reflectance spectra as preceived colors. We have implemented Endler's (REF) segment analysis model, which approximates visual models but does not directly use sensory information; the model of Osorio & Vorobyev (REF), which provides a flexible framework for visual modeling; and the tetrahedral color space (GOLDSMITH, ENDLER, STODDARD) which has been extensively developed to represent colors in the avian vision color space.

4.2 Spectral Shape Analysis

4.2.1 Trichromatic variables

Obtaining trichromatic color variables (related to hue, saturation and value) is pretty straightforward in **pavo**. Since reflectance spectra is stored in an object of class **rspec**, the **summary** function recognizes the object as such and extracts 23 variables, as outlined in MONTGOMERIE. Though outlined in a book chapter on bird coloration, these variables are broadly applicable to any reflectance data, particularly if the taxon of interest has color vision within the UV-human visible range.

The description and formulas for these variables can be found in Table 1.

```
> summary(sppspec)
```

	B1	B2	B3	S1.UV	S1.violet
cardinal	8983.870	22.40367	52.96101	0.16618070	0.19132718
jacana	9668.442	24.11083	53.48957	0.09623268	0.11048401
oriole	9107.568	22.71214	54.31982	0.07474281	0.08314578
parakeet	6022.866	15.01962	28.82353	0.16803871	0.18515359
robin	5741.033	14.31679	37.96515	0.08500350	0.10012432
tanager	8516.756	21.23879	29.76561	0.20342675	0.23931496

	S1.blue	S1.green	S1.yellow	S1.red	S2
cardinal	0.11784952	0.1900225	0.2517886	0.5332009	7.329535
jacana	0.19475522	0.3082406	0.2479744	0.4079294	6.883918
oriole	0.05721376	0.3256520	0.3630756	0.5491879	12.992914
parakeet	0.14334383	0.4252528	0.3397279	0.2717663	4.768352
robin	0.14429801	0.2675057	0.2668015	0.5099805	9.172749
tanager	0.26109724	0.3199401	0.2430633	0.2238215	3.120773

	S3	S4	S5	S6	S7
cardinal	0.3682957	0.15813546	4004.193	45.73531	-0.38403228
jacana	0.2469875	0.01690805	3312.697	45.71935	-0.46942827
oriole	0.3444943	0.06767440	5232.176	50.13910	-0.70982859
parakeet	0.4467452	0.23727515	1900.192	22.77878	-0.20889035
robin	0.3031191	0.01271259	2531.901	33.82625	-0.57712488
tanager	0.3341862	0.15240014	1105.252	20.22771	-0.08198454

	S8	S9	S10	H1	H2	H3
cardinal	2.0414210	-0.8373693	12.909318	700	425	596
jacana	1.8962165	-0.7357544	112.148743	700	388	504
oriole	2.2075903	-0.9257720	32.620756	700	388	500
parakeet	1.5166018	-0.5816948	6.391743	568	613	502
robin	2.3626975	-0.8148189	185.854859	700	349	500
tanager	0.9523946	0.0359801	6.249303	557	600	428

	H4	H5
cardinal	1.54995089	617
jacana	0.85480076	463
oriole	1.15606660	538
parakeet	0.59271149	500
robin	1.15020735	637
tanager	0.04509807	512

Color		
Variable	Equation	Description
B1	$\sum_{\lambda=300}^{700} R_{\lambda}$	Total brightness, total reflectance
B2	B_1/n_{wl}	Mean brightness.
B3	R_{\max}	Intensity.
S1		Chroma, spectral purity.
S2	R_{\max}/R_{\min}	Spectral saturation
S3		
S4		
S5		
S6		
S7		
S8		
S9		
S10		
H1	$\lambda_{R_{\max}}$	Hue: wavelength of peak reflectance
H2		
H3		
H4		
H5		

Table 1: The complete set of tristimulus variables calculated by `summary` in `pavo`

4.2.2 Peak shape descriptors

Particularly in cases of reflectance spectra that have a single, discrete peak, it might be useful to obtain variables that describe that peak's properties. The `peakshape` function identifies the peak location (`H1`), returns the reflectance at that point (`B3`), and identifies the wavelengths at which the reflectance is half that at the peak, calculating the wavelength bandwidth of that interval (the **Full Width at Half Maximum**, or `FWHM`). The function also returns the half widths, which are useful when the peaks are located near the edge of the measurement limit and half maximum reflectance can only be reliably estimated from one of its sides.

If this all sounds too esoteric, fear not: `peakshape` has the option of returning plots indicating what it's calculating. The vertical continuous red line indicates the peak location, the horizontal continuous red line indicates the half-maximum reflectance, and the distance between the dashed lines is the `FWHM`:

```
> par(mfrow=c(2,3),mar = c(5, 4, 0.5, 0.5) + 0.1)
> peakshape(sppspect, plot=T)
```

	B3	H1	FWHM	HWHM.l	HWHM.r	status
cardinal	52.89508	687	104	100	4	OK
jacana	53.90258	700	NA	93	NA	OK
oriole	54.40384	692	148	141	7	OK
parakeet	30.05703	575	125	66	59	OK
robin	37.91676	700	NA	107	NA	OK
tanager	30.71021	555	232	142	90	OK

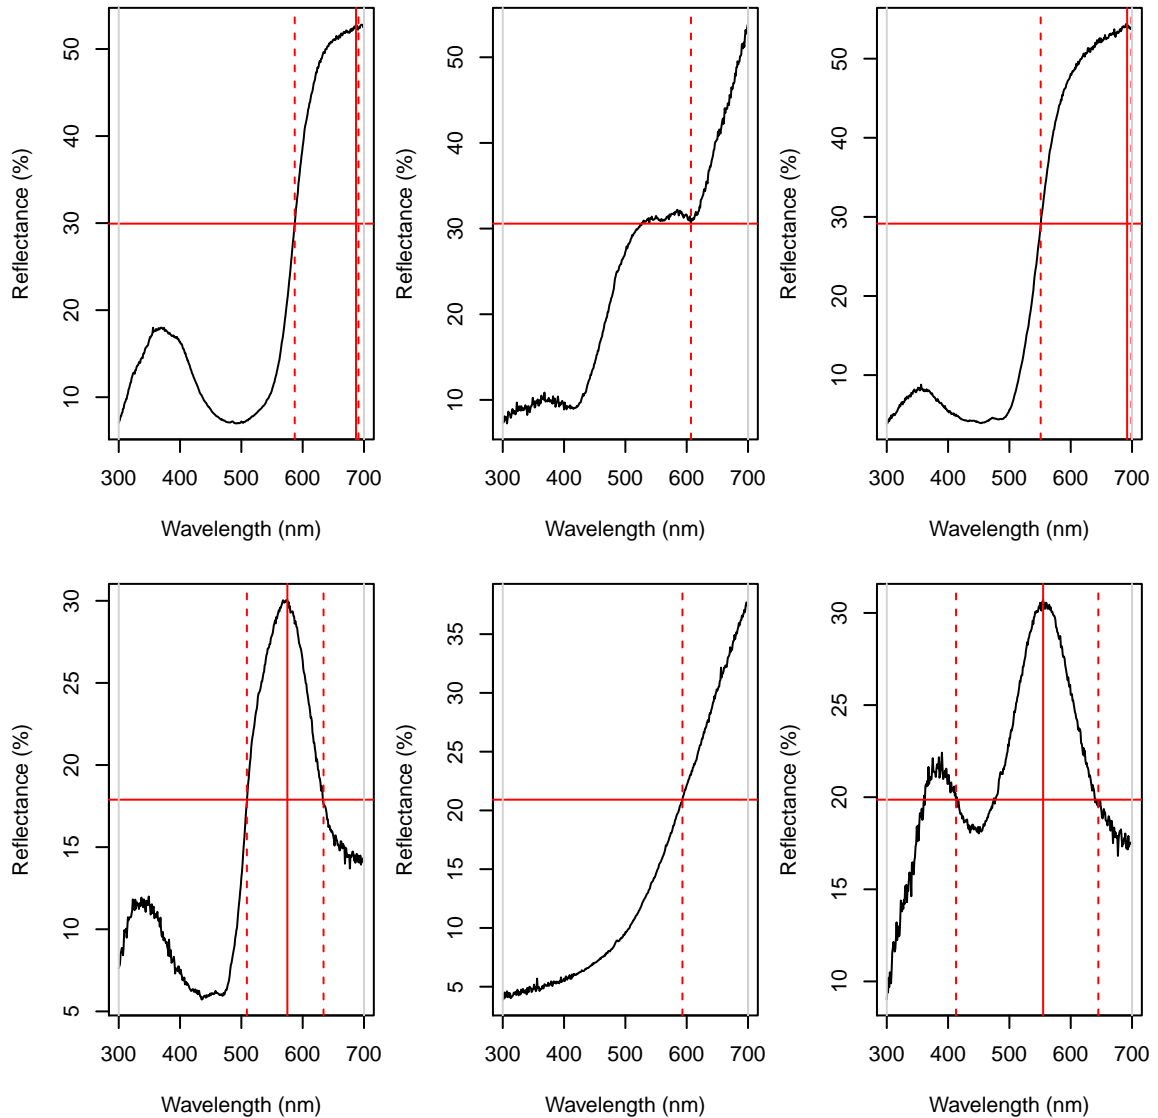


Figure 4: Plots from `peakshape`

As you can see, the variable is meaningless if the curve doesn't have a clear peak. Sometimes, such as in the case of the Cardinal (Figure 4, first panel), there might be a peak which is not the point of maximum reflectance of the entire spectral curve. The half-width can also be erroneously calculated when there are two peaks, as can be seen in the case of the Tanager (Figure 4, last panel). In this case, it's useful to set bounds when calculating the FWHM, using the `bounds` argument. `peakshape` also offers a `select` argument to facilitate subsetting the spectra data frame:

```
> peakshape(sppspec, select=2, bounds=c(300,500), plot=T)
```

	B3	H1	FWHM	HWHM.l	HWHM.r	status
cardinal	17.99853	369	100	47	53	OK




Figure 5: Plot from `peakshape`, setting the bounds to 300-500nm

Examples

```
> hist(rnorm(50))
```

More examples

Some more examples: